

A Bug Reproduction Approach Based on Directed Model Checking and Crash Traces

Mathieu Nayrolles¹, Wahab Hamou-Lhadj¹, Sofiène Tahar² and Alf Larsson³

¹Software Behaviour Analysis (SBA) Research Lab, ECE, Concordia, Montréal, Canada

²Hardware Verification Group (HVG) Research Lab, ECE, Concordia, Montréal, Canada

³PLF System Management, R&D Ericsson, Stockholm, Sweden

mathieu.nayrolles@gmail.com, wahab.hamou-lhadj@concordia.ca, tahar@ece.concordia.ca, alf.larsson@ericsson.com

May 2nd, 2016

Context: Software are released with bugs.

- Despite testing and verification, softwares are pledged to be released with latent bugs.
- Latent bugs will cause field crashes / faillures.
- Patching field faillures is challenging:
 - We have to know about them.
 - Information is scarce and inconsistent
 - Most valuable information are the one that help to reproduce a bug [Bettenburg, 2008].

Related Works: Current ways to reproduce a crash.

Record and replay

- Instrumentation of source code
- Record on-field execution
- Replay in-house
- Cheap & easy to implement
- Yield good results
- Overhead (1% to 1066%) and privacy concerns
- JRapture'00, BugNet'05, ReCrash'08

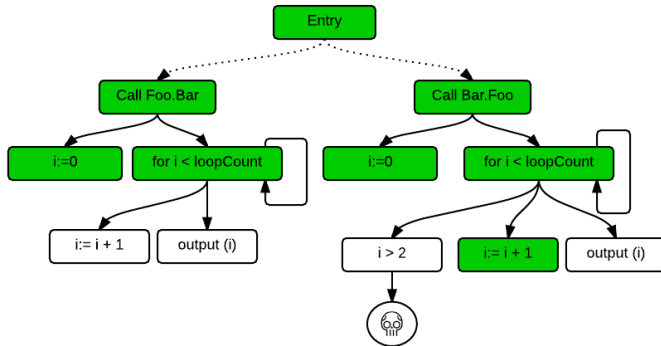
In house crash reproduction

- Core dump.
- Forward Symbolic Execution.
- Backward Symbolic Execution.
- Yield average results.
- NP-Complex Problem
- Exponential learning curve.
- Privacy concerns.
- BugRedux'12, RECORE'13, STAR'13

JCHARMING: A different direction

- Avoid code instrumentation
 - 0% overhead
- Do no yield privacy concerns
 - Scalable to real-world and industrial/proprietary software systems
- Leverage the stack traces resulting from a crash
 - More and more often present in bug reports
- JCHARMING uses directed model checking and backward slicing.


Preliminaries: **Testing**, Model Checking and Directed Model Checking



- Depends on the tester understanding of the SUT.
- Inefficient because it is not exhaustive.

Preliminaries: Model Checking

- Checks if a given system under test (SUT) meets a specification p by exhaustively testing every states. [Visser, 2003], [Kropf, 1999]




- Ensures that p is reached at some point and not that p holds nor $\forall x, p$ is satisfiable.
- We aim to verify that \forall states the program does not crash:

$$\forall x. (SUT, x) \models \neg c$$

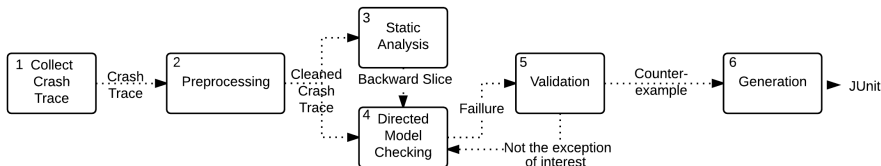
- Explores each and every state of the program, hence it is complete.
- Impractical for real-world and large systems

Preliminaries: Testing, Model Checking and **Directed Model Checking**



- Explores only the states that may lead to a specific location. [Rungta, 2009]
- Use insights — generally heuristics — about the SUT to prune states.

The JCHARMING approach



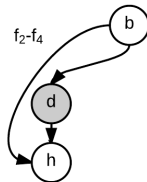
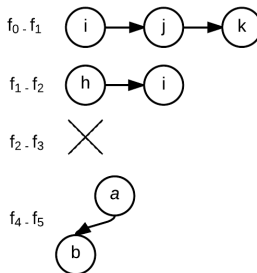
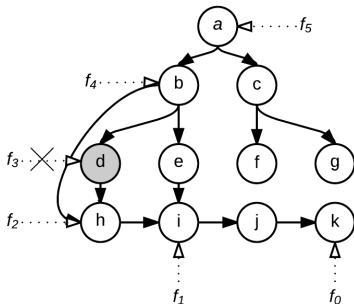
• Step 1: Collect the crash trace

```
1. javax.activity.IAE:loopTimes should be < 3
2.  at Foo.bar(Foo.java:10)
3.  at GUI.buttonActionPerformed(GUI.java:88)
4.  at GUI.access$0(GUI.java:85)
5.  at GUI$1.actionPerformed(GUI.java:57)
6.  caused by java.lang.IndexOutOfBoundsException : 3
7.  at saner.Foo.buggy(Foo.java:17)
8.  and 4 more ...
```

• Step 2: Preprocessing

Step 3: Building the Backward Static Slice (Cont'd)

- A backward slice contains all possible branches that may lead to a point n from a point m as well as the definition of the variables that control these branches
- Let's assume we have $T = \{k, i, f, d, b, c\}$
- Without using T the backward slice is $\{a, b, d, e, h, u, j, k\}$.
- With T , and frame by frame it is $\{a, b, d, h, i, j, k\}$.



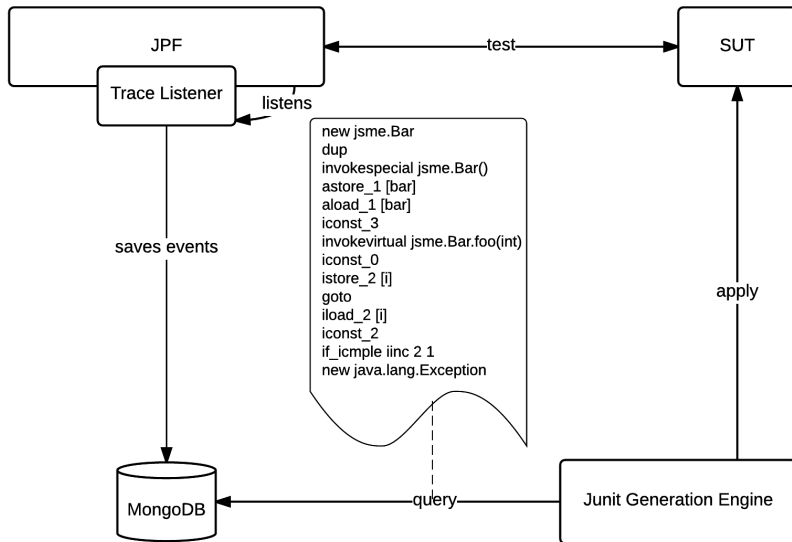
Step 4: Directed Model Checking

- We use Java PathFinder (JPF)
 - JVM for Java bytecode verification.
 - Front-end for the SPIN model checker.
 - Developed and maintained by NASA.
- Generate states
- Forward
 - Generates the next state S_{t+1} and add it to the backtrack table
- Backward
- Backtrack
 - Restore the last state of the backtrack table.
- Restore state
- Check properties
 - Is triggered after each forward, backward and restore operations.

Step 4: Directed Model Checking (Cont'd)

- We modified the *generate states* and the *forward steps*.
- The *generate states* is populated with: $\bigcup_{i=0}^{entry} bslice_{[f_{i+1} \leftarrow f_i]} \subset SUT$
- The *forward step* can explore a state if s_{i+1} and the transition x from s_i to s_{i+1} are in $\bigcup_{i=0}^{entry} bslice_{[f_{i+1} \leftarrow f_i]} \subset SUT$
- The *check properties* step of JPF:
 - If the current states transitions x yield an exception
 - We execute x and compare the stack trace to the original
 - If the two exceptions match, the bug is *reproduced*.
- JPF is now directed and explores only a sub-system of the SUT.

Step 6: Generating Test Cases for Bug Reproduction



- 30 bugs belonging to 10 open sources systems (Ant, ArgoUML, DnsJava, jFreeChart, Log4j, MCT, PDFBox, Mahout, Hadoop, ActiveMQ).
- 80% success ratio (24/30).
- Average success time is 19 minutes.
- Average fail time is 11 minutes.
- Average JUNIT length is 5 java statements.
- Only 6% (2/30) of the bugs could have been reproduced using undirected model checking.

Conclusion

- JCHARMING (Java CrasH Automatic Reproduction by directed Model checking)
- Automatic bug reproduction technique that combines crash traces and directed model checking
- Direct the model checking engine with a backward static slice
- Was able to reproduce fully or partially 80% of the bugs
- Stress JCHARMING with more bugs
 - Fine tune the approach
 - Assess the scalability on larger / proprietary systems
- Test the performances of JCHARMING with multi-threading related bugs

QUESTIONS?